

Sample Book

Author

May 11, 2022

Abstract

Here I write the abstract

Contents

Abstract	iii
1 Introduction	1
2 State of the Art	3
2.1 The problem with evidence	3
2.2 Reasoning with evidence	3
2.3 Bayes!	4
2.4 Problems with the Bayesian Approach	4
3 Simulations to Evaluate Bayesian Networks.	7
3.1 Introduction	7
3.2 Setting-up an Experiment	7
3.2.1 Simulation	8
3.2.2 Reporters	9
3.3 Creating a Bayesian Network from a Simulation Automatically	10
3.4 Evaluating the Bayesian Network	11
3.4.1 Structural Criteria	11
3.4.2 Performance Criteria	11
3.4.3 Human Criteria	11
4 Simple Non-Spatial Simulations	13
4.1 Introduction	13
4.2 Method	14
4.2.1 Behavioural rules for the simulation	14
4.3 Results	15
4.4 Discussion	15
4.4.1 Structural Criteria	15
4.4.2 Performance Criteria	19
4.4.3 Human Criteria	21

4.4.4	Comparison to Vlek 2015	21
4.4.5	Quality of scenario	22
4.5	Important take-aways	23
4.5.1	We don't need the scenario node for mutual exclusivity	23
4.5.2	Resolution and number of runs	23
5	Simple Spatial Simulations	25
5.1	Introduction	25
6	The Great Grote Markt Robbery Island Prior	27
6.1	Introduction	27
7	Conclusion	29
7.1	Conclusion	29
7.2	Future Work	29

Chapter 1

Introduction

Reasoning with evidence in law is cool!

The main research question for this Master's Thesis is: can we automatically create good Bayesian Networks that reflect the ground truth of simulations? If so, can we use this simulation + Bayesian Network setup to investigate BN idioms and methods for law more generally, to see how well the probabilistic approach holds up?

Chapter 2

State of the Art

2.1 The problem with evidence

When we find evidence for a hypothesis that we have held in the back of our minds, we then find the hypothesis more likely. This is the basic idea behind reasoning with evidence. In a constellation of hypotheses and pieces of evidence, we want to construct a network that will lead us to believe as many true hypotheses as possible, given the evidence that we have.

However, evidence itself is elusive, and it's connection to hypotheses is as well - how can we be sure that some evidence supports some hypothesis, and even if we know that it does, how can we express how strong the piece of evidence is? Some evidence is very weak, and only after a tedious process of ruling out other factors and careful investigation and collection of other pieces of evidence, we can come to a conclusion about a hypothesis. On the other hand, some evidence is so strong that it leaves no room for doubt.

We all have intuitions about evidence strength - but can we make these intuitions precise? Additionally, can we manage the complex realities of weak evidence for many different hypothesis? These are questions that have guided this state of the art section.

In this section I will briefly discuss methods for reasoning with evidence, then transition to probabilistic reasoning with evidence, Bayesian Networks. I will discuss some problems with Bayesian Networks.

2.2 Reasoning with evidence

Argumentation graphs and their semantics. Scenario theory. Bayes Law. Hybrid Approach.

2.3 Bayes!

Probabilistic reasoning (Dahlman, simple). Bayesian Networks Fenton and Vlek - combining scenarios, idioms in Bayesian Networks.

2.4 Problems with the Bayesian Approach

There are many problems with the Bayesian Approach.

The most obvious is the problem of the numbers: where do we get them? We can get some of them from statistics, but we need to many numbers that we have to make some numbers up. This is not necessarily a bad thing - we put our (betting) money where our mouth is and assign numerical precise probabilities to situations that we preciously only had vague intuitions for. However, this brings about the veil of objectivity. By giving a probability to your intuitions, you have made your intuitions more precise and you can now reason with them, update on evidence using Bayes Law, and everything's great. In some domains, this is obviously okay. If you want to bet cents on world events and walk that fine line between calibration and discrimination for fun and profit, that's no problem. After all, there are incentives to abstain from the unclear, the stuff that might not have a specified answer, the vague. But when we are talking about using evidence in law, we are talking about exactly that domain - we're not making predictions about the price of oil in 6 months, or the outcome of the French election, with clear outcomes, clear procedures for measurement. Instead, we're trying to make predictions about crimes and crime scenarios, which are a lot vaguer, and strangely unobservable at times - things like motives, or behaviours that happen under specific circumstances, interlocking stuff with complex dependencies. We all have intuitions about evidence strengths in vague situations, but they are more difficult to make precise than the traditional 'forecasting' events. So trying to assign probabilities without a clear method of calibration, makes that they will be imprecise.

Level of granularity.

Independence relations. Selecting events.

Technical problems with Bayesian Networks. Precision. Evaluation. Sensitivity Analysis.

1. The problem of the reference class, which is also a problem within a 'clear' domain - but this is a fundamental philosophical problem that I will not discuss further.
2. We do not know the probabilities (frequency) of our variables in the first place.
3. We do not know how robust the network is to imprecise or wrong frequencies.
4. We do not know if the assumption of independence between any two variables holds. The assumption of independence is necessary in Bayesian Networks, otherwise the

complexity of the network becomes unmanageable.

Chapter 3

Creating Simulations to Evaluate Bayesian Networks.

3.1 Introduction

In this part, I explain the general method for creating simulations with automatic Bayesian Networks, as well as how to evaluate these Bayesian Networks. The general process of creating simulations to evaluate Bayesian Networks is illustrated in Figure 3.1. We start by defining a simulation with agents, and then have reporters report on that simulation in the ‘Experiment’ stage. Relevant events in simulations are collected by the reporters in each run. The collection of runs is used by an automatic Bayesian Network constructor algorithm to construct a Bayesian Network automatically in the ‘Building BN’ stage. Finally, the constructed Bayesian Network is evaluated with respect to the criteria in the ‘Evaluate BN’ stage. These three stages are explained in this section. Specific instances of simulations and networks I created with this process are the subject of the next chapters.

3.2 Setting-up an Experiment

An experiment consists of a simulation and reporters. Reporters are defined separately from the simulation because they are not inherent to it - they are defined with respect to what we want to know about the simulation.




Figure 3.1: Method for evaluating automatically generated Bayesian Networks from simulations.

3.2.1 Simulation

Structure of a simulation

In the simulation, we are simulating some sort of criminal scenario - a theft, usually. Or we are simulating purely for the theoretical things. The simulation can be as precise as necessary, but there are certain things that need to be present: we need to have states that happen, there needs to be evidence for those states as well. The granularity of the simulation and its complexity depend on the modeller and her requirements.

Spatial and Non-spatial simulations.

I'm discussing two types of simulations: spatial simulations and non-spatial simulations. In a spatial simulation, an agent's behaviour is mediated with respect to their environment - eg, agents cannot pass through buildings, or cannot see agents that are far away, or can only steal from another agent when that agent is nearby. In a non-spatial simulations, agents can behave and interact with each other, but this is happening without any environment, hence we are simulating an abstraction. In concrete terms, you can think about a non-spatial

simulation as a communication game.

For this project, that means that spatial simulations are more complex and interesting than non-spatial simulations, as there are more possibilities for variety.

The simulations were programmed in MESA, a python package that is made for the creation of Multi-agent systems simulations. We define an environment that the agents can interact with, as well as agents that perform some behaviours. Specifics of simulations and agent behaviour are described in the following chapters.

Predictability and randomness

Where does the interest of the simulation come from? In one part, agents sometimes do things because they are commanded to do so by the computer ¹. This means that, at the start, a random number generator might ‘decide’ that some agent has a motive, since the random number generator generated a 1 instead of a 0. On the other hand, some randomness arises from interactions between agents, or between agents and their environments. This is where spatial simulations bring additional value compared to non-spatial simulations.

In non-spatial simulations, all agent states are essentially brought about by a combination of randomly generated numbers, and reasoning rules. For example, if an agent has a tendency to lie (randomly generated), and it has the opportunity for lying (brought about due to the current non-spatial simulation), then it will lie. Hence a combination of behavioural rules and randomly generated numbers results in a state of ‘agent lies’ of 1.

However, in spatial simulations, interactions and behavioural rules and randomly generated numbers are all brought together: if an agent is near another agent (chance interaction), and it has a tendency to steal (randomly generated), then it will attempt (but might not succeed) in stealing. Here the behavioural rule might lead to a more interesting/complex/complicated outcome than in the non-spatial simulation.

3.2.2 Reporters

In the simulation, certain states can be brought about. For example, an agent can succeed in stealing an object, or in lying, or in having a motivation (or in not having those things). We need a way to observe these states: this is where reporters come in. A reporter reports the outcome of a relevant event or state in the simulation, and is embedded in the code. If an event happens (or does not happen), the reporter reports that the event is true (or false). In essence, the reporter (R) is a random variable (RV) that maps an event (e) to a truth value:

¹rephrase, this is always true lol.

$$R : e \rightarrow \{0, 1\}$$

Not all the states in the simulation have a reporter associated with it - otherwise I could build infinitely many reporters. I could have reporters for names, for *agent_Q_at_x_1_y_200*. Hence, I only created reporters for states I deemed relevant for the scenario that I am investigating. Here is a subjectivity gap. I can imagine that in my simulation of the Grote Markt (see later chapter), there is some part of the simulation that by chance geometry, lends itself to an easier job for the thief than another part of the simulation. If the thief and victim spawn near this point, then the probability of the thief succeeding will be higher. Increasing the granularity of the reporters might help us determine if there is a spot like this. However, I did not implement this level of granularity in the simulation (yet), because that is a local and specific part, and does not fit into the more global scenario description (the scenario of theft is spatially-free).

The global state of one run of the simulation, is the combination of all reporters.

$$G = (e_0 \rightarrow \{0, 1\} \times e_1 \rightarrow \{0, 1\} \times \dots \times e_n \rightarrow \{0, 1\})$$

or,

$$G = R_1 \times R_2 \times \dots \times R_n$$

for n reporters.

Then, we collect these global states over the number of runs that we do for each experiment, which results finally that the output O of this stage of the method, is a series of global states, one for each run:

$$O = (G_0, G_1, \dots, G_{runs})$$

3.3 Creating a Bayesian Network from a Simulation Automatically

The output of an experiment is the collection of runs O , where each run is the global state G of the simulation, as measured by the random variables R . Semantically, it fits that reporters are random variables, as the reporters become the nodes in the Bayesian Network.

Once we have a collection of states and runs, we can give this to an automated Bayesian Network learner, such as those implemented in pyAgrum. These learners can interpolate

a Bayesian Network using algorithms, such as Greedy Hillclimbing and K2. This results in automatic generation of the Bayesian Network which is solely based on the data that is collected in O .

For more detail on these algorithms watch this space.

3.4 Evaluating the Bayesian Network

We want to evaluate different aspects of the Bayesian Network: structural criteria, performance criteria and human criteria.

3.4.1 Structural Criteria

1. (temp) Events are ordered temporally - scenario-like.
2. (con) Evidence connects to hypotheses.
3. (exc) All events that are irrelevant are not included in the scenario BN.
4. (exh) All events that are relevant are included in the scenario BN.
5. (ind) Independent events are not connected to each other.

3.4.2 Performance Criteria

1. The accuracy of the network is not lower than the inherent randomness of the simulation.
2. The strong view: probabilities in the network correspond exactly to probabilities in the simulation.
3. The weak view: Updates on evidence in the network go in the same direction as updates on evidence in the simulation.
4. Sensitivity analysis shows that no simulation-irrelevant event influences some output ²

3.4.3 Human Criteria

1. Do we think that a human can find these numbers?
 - (a) How robust is the network against disturbances around the mean (problem with precision)?
 - (b) How robust is the network around rounding to arbitrary intervals?

²rephrase

Now that we've established how the automatically generated Bayesian Networks are going to be judged, we can show cautiously in the next chapter how well they are holding up!

Chapter 4

A Simple Stabbing - Non-Spatial Simulations and their Bayesian Networks

4.1 Introduction

Here I talk about the credibility game and Charlotte Vlek's Stabbing Bayesian Network example.

In these non-spatial simulations, there are agents, and they interact with each other, but there is no environment for them to interact in. So the simulation is a pure combination of the probabilities assigned to each state transition (or something along those lines). In a sense, the probability for an agent to 'stab with knife' purely depends on the probability of 'motive' and 'opportunity'. Compared to a spatial simulation, where 'opportunity' is more complex, as it involves proximity to victim which arises from the simulation and not from some random number generator.

These non-spatial simulations are boring, but they are necessary first steps: after all, if we cannot make BNs out of these predetermined (the probabilities in each run are not predetermined, but the distributions that they are drawn from, and their thresholds, are), then the rest of our endeavours will be fruitless. On the other hand, if the process for creating and evaluating these simple simulations work well, then we can proceed to modelling more complex, spatial simulations.

In this part, I have two experiments. One simple experiment meant to be a replication of Charlotte Vlek's Bayesian Network in xxx ¹, and another simple experiment mainly

¹rephrase

meant to test the evaluation criteria for the Bayesian Networks, as outlined in the previous chapter.

4.2 Method

Take the Vlek networks from Vlek Jurix 2015.

The general story is: Jane and Mark had a fight, but Jane had a knife. Mark died.

Then, there are two specific scenarios that can explain why Mark died. In scenario one, Jane stabbed Mark, and then he died. In scenario 2, Jane threatened Mark with the knife, Mark hit Jane, Jane dropped the knife, Mark fell on the knife, and Mark died by accident. There are two separate networks for these scenarios.

I'm going to create two separate networks, and then also see if I can merge them, by creating a Jane-and-Mark-knife simulation, assigning some random probabilities that correspond to the story, and see what the K2 algorithm makes of it. Then I will also merge the two networks to see if the K2 can deal with mutually exclusive nodes (eg: 'Mark died by accident' should rule out 'Mark died by stabbing').

So, I created some logical rules, either atoms or rules, and we can process in the simulation these using standard forward chaining inference. Every atom has a prior probability, and every conclusion of a rule has a probability given the F/T state of the premises. At every step, the simulation checks which new sentences are true, applies a rule with a given probability, and counts the outcome.

This does mean that rules at the end are not triggered as often as rules in the beginning, even if they have the same trigger probability (they are on other sides of the chain, more needs to be have happened to conclude that Mark died).

4.2.1 Behavioural rules for the simulation

This was done using forward chaining, with a time index, which means that a rule could only be triggered at one time (otherwise the probabilities get messed up). So we have a forward chaining rule, which means that if the premises of some rule are true, there are no excluding facts true, the timestep is correct, and the random number generator generated a number that is lower than the probability threshold, we find that the conclusion is true, and add it to our found facts. We keep doing this.

Vlek's paper contains no probabilities, so I'm just making some up. The logical sentences are reporters in the simulation, and then I let the simulation run for 10,000 times.

Premise	Conclusion	P(conclusion) given premises
	Jane and Mark fight	20
	Jane has knife	70
Jane and Mark fight, Jane has a knife	Jane stabs Mark with knife	1
Jane and Mark fight, Jane has a knife	Jane threatens Mark with knife	3
Jane threatens Mark with knife	Mark hits Jane	90
Mark hits Jane	Jane drops knife	50
Jane drops knife	Mark falls on knife	10
Mark falls on knife	Mark dies by accident	60
Jane stabs Mark with knife	Mark dies	70
Mark dies by accident	Mark dies	100

Table 4.1: For combined scenarios: blue rules belong to scenario 1, red rules belong to scenario 2, and black rules belong to both.

4.3 Results

We succeeded in generating three different BNs: One for scenario 1, one for scenario 2, and one for the combined network (Figure 4.1, Figure 4.3, Figure 4.5).

4.4 Discussion

How well do these networks evaluate compared to the criteria we set out initially, and how well do they compare to the Vlek networks, and the criteria set out in Vlek 2015?

4.4.1 Structural Criteria

1. (temp) Events are ordered temporally - scenario-like.

In the forward chaining knowledge base, we have premises and conclusions. If we take the event in a conclusion, then we know that the premise of that rule, is the parent of that node in the Bayesian Network. (again, rephrase this is so vague. Maybe make it a rule or smth). The two “premiseless” events Jane and Mark fight, and Jane has a knife, are also parentless in the simulation. In this sense, the structure of the Bayesian Network reflects the structure of the inference rules that we set out.

We see that in both the separate networks, the temporal ordering of the nodes is correct - every node has as a parent a node that represents an event that would have happened beforehand. In the combined network, this is not the case: mostly the temporal links are represented ok (eg: the chain of events from ”Jane threatening Mark” to ”Mark dies by accident” is represented correctly), however, this chain of nodes all has as a parent ”Mark dies”, which would, temporally, occur only after




Figure 4.1: Automatically generated BN with K2 and the above forward chaining rules, scenario 1.




Figure 3. A network structure for the scenario about Jane stabbing Mark

Figure 4.2: Vlek BN.




Figure 4.3: Automatic generated BN with K2 and above forward chaining rules, scenario 2

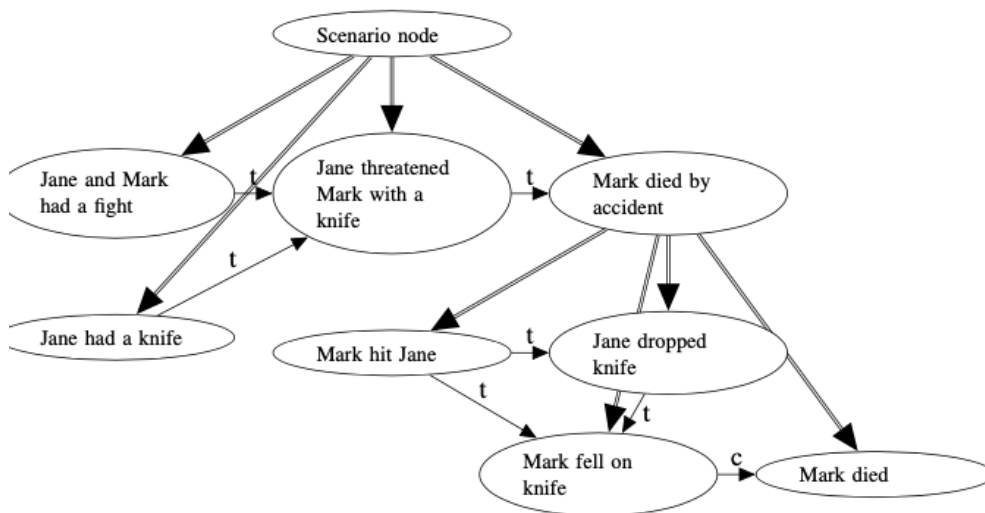


Figure 4. A Bayesian network structure for the second scenario

Figure 4.4: Vlek BN.

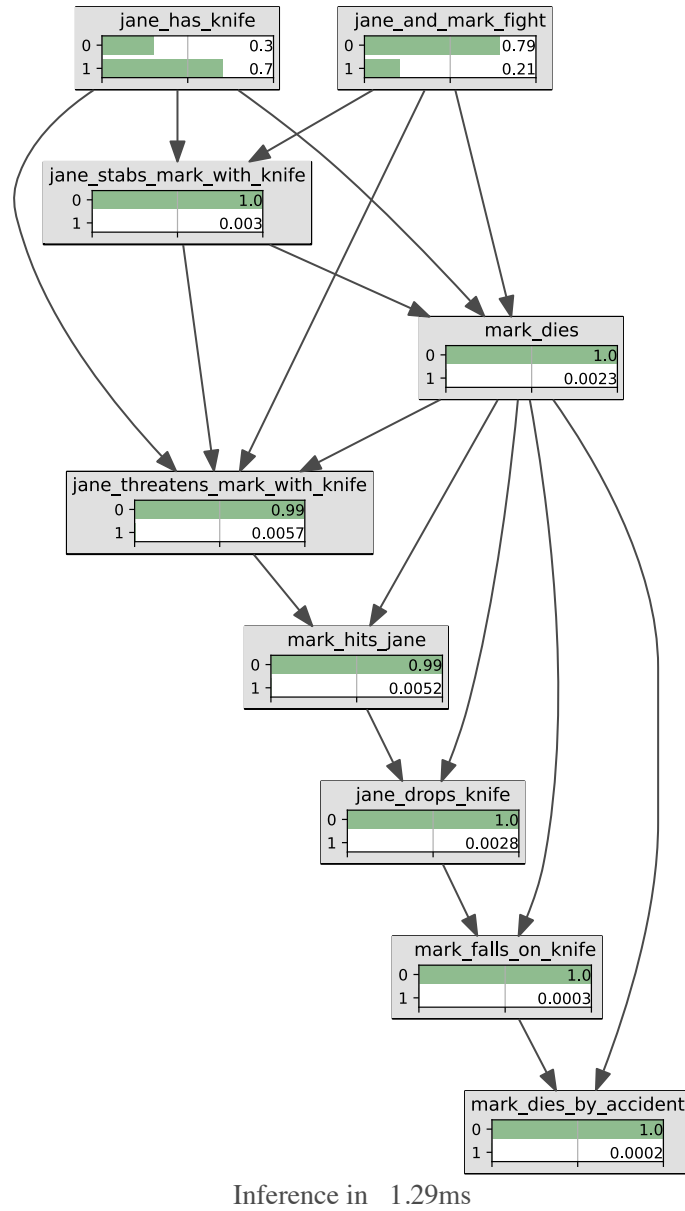


Figure 4.5: Automatically generated BN with rules from both scenarios included.

"Mark falls on knife" occurs. This means that in the combined Bayesian Network, we cannot satisfy the temporal ordering constraint due to the conditional probabilities between the events (eg: if Jane stabs Mark, then she cannot threaten him after, which is a relation that 'inhibits' the 'Jane threatening Mark' node - eg: there is a conditional relation between the two nodes, but not a temporal or causal one). As a conclusion about this criteria: Maybe temporal ordering should just hold within scenarios and not between scenarios, because we get into these sorts of causal messes. Or we should use the scenario node to organise different scenarios together. Idea: scenario node could cause temporal ordering, but what in cases with much shared evidence?

2. (con) Evidence connects to hypotheses.

In the original network, we have no evidence, so it is not included into this one. However, I might add some evidence to test the Performance Criteria.

3. (exc) All events that are irrelevant are not included in the scenario BN.

This criteria, and the next one, are freebies, because we are just copying a network that already exist - our set of relevant events is the same as that in the original network by Vlek. We do not have any irrelevant event because we can see that all events are connected to each other, and all events reflect an underlying 'decision' by an agent.

4. (exh) All events that are relevant are included in the scenario BN.

We have included all relevant events - all the events in the knowledge base are also in the network.

5. (ind) Independent events are not connected to each other.

This is not entirely correct in the uncombined networks (and also not correct in the combined network) - If Mark dies by accident, this is only dependent on Mark falling on the knife, and not also on Jane dropping the knife (since all cases where Mark dies by accident are also cases where Mark falls on the knife, we don't need a separate condition for Jane dropping the knife). This is probably due to an insufficient number of runs (eg: the probability that this happens is very low and we might not have enough data to correctly estimate it).

4.4.2 Performance Criteria

1. The accuracy of the network is not lower than the inherent randomness of the simulation.
2. The strong and the weak view.

Conclusion	P(conclusion) given premises	P(event) given premises
Jane and Mark fight	20	20.87
Jane has knife	70	70.23
Jane stabs Mark with knife	1	2.05
Mark dies	70	75.46

Table 4.2: For only scenario 1.

Conclusion	P(conclusion) given premises	P(event) given premises
Jane and Mark fight	20	20.87
Jane has knife	70	70.23
Jane threatens Mark with knife	3	3.88
Mark hits Jane	90	90.21
Jane drops knife	50	50.09
Mark falls on knife	10	11.13
Mark dies by accident	60	61.25
Mark dies	100	99.79

Table 4.3: For only scenario 2

The strong view: probabilities in the network correspond exactly to probabilities in the simulation. The weak view: Updates on evidence in the network go in the same direction as updates on evidence in the simulation.

For both scenarios, the probabilities in the network are very similar to the probabilities in the ground truth (simulation). There is a loss of precision, which is the worst in case of ‘Mark dies’ in the first scenario, which is estimated at 75% by the network, but in fact in the simulation is only 70% probable. I don’t know where this comes from. Apart from that, all divergences for the original probabilities are within 2%. This means the strong view does not hold, but the weak view does (updating on evidence gets us with probabilities that are closer to the simulation probabilities than before).

The probabilities of the conclusions in the combined network given the premises, also look relatively similar on a human scale - I assume that a human probability elicitor would not be able to distinguish between a probability of 11.13, 10.92 or 10.00 for the conclusion “Mark falls on knife”. On the other hand, if we round all the probabilities in the network to 10, a human probability elicitor might be able to say that the probability of ‘Mark falls on knife’ is 0.1 (or 10%), to exactly that level of precision (so not 0.10, or 0.100). The resulting probabilities might be elicitable, while still meaningfully reflecting the ground truth.

3. Sensitivity analysis shows that no simulation-irrelevant event influences some output

Conclusion	P(conclusion) given premises	P(event) given premises
Jane and Mark fight	20	20.87
Jane has knife	70	70.17
Jane stabs Mark with knife	1	2.06
Jane threatens Mark with knife	3	3.88
Mark hits Jane	90	91.50
Jane drops knife	50	52.97
Mark falls on knife	10	10.92
Mark dies by accident	60	66.33
Mark dies (premise: Jane stabs Mark)	70	68.68
Mark dies (premise: Mark dies by accident)	100	98.77

Table 4.4: For combined scenarios

2

4.4.3 Human Criteria

1. Do we think that a human can find these numbers?

The conditional probability tables as displayed in these networks are not great.

- (a) How robust is the network against disturbances around the mean (problem with precision)?
- (b) How robust is the network around rounding to arbitrary intervals?

4.4.4 Comparison to Vlek 2015

If we only look at the ordering of the Bayesian Network, we can see that for Figure 4.1, the sub-scenario structure is the same as in Vlek 2015 Figure 3, our Figure 4.2. There's no scenario node constraining the network, all the information is contained in the network, no scenario node needed (todo: make the nodes the exact same names).

However, there are differences between Vlek 2015 Figure 4, and the automatically generated BN here. Figure 4 in Vlek 2015 is replicated below in Figure 4.4 (todo: ask permission?? or just remove). The automatically generated BN is very linear and can be interpreted in a purely temporal way: mark hits jane, jane drops the knife, and due to jane dropping the knife, and mark falling on it, mark dies by accident - and if mark dies by accident, mark dies. The probabilities for all these events (from jane dropping the knife on), are ridiculously low, and don't really make sense (should interpret the small probabilities as e, and not as actual numbers I guess, due to underflow?).

²rephrase

In Vlek’s paper, we see a subscenario: we have the subscenario “Mark died by accident”, which contains events such like mark hit jane, which leads to jane dropping the knife, and mark falling on the knife, and then dying. The coparents of Mark died are the same in this network as in the automatically generated BN, however, we once-again miss the scenario-like construction where mark hitting jane, jane dropping the knife, and mark falling on the knife are connected as part of a subscenario, rather than their “own” nodes.

Why to choose for subscenarios when it is not required? Tomorrow I will look up why the scenario construction was used and I will see if it does something that I miss? Coherence? But that also travels up the chain? But something with d-separation probably. I’ll read the Vlek 2015 again i guess.

In the full final network, we see that the two mutually exclusive scenarios (stabbing vs falling on knife) are excluding each other even without the use of a scenario node - so we can make networks that combine two different stories. This is nice.

4.4.5 Quality of scenario

- Completeness - have all the parts and not one more. Not a big deal for this simulation, as I’ve taken all the nodes from Vlek 2015, there’s no discussion about granularity.
- Consistent - no internal contradictions. There is no explicit structure in the final BN that ensures consistency. The probabilities inside the BN need to ensure that. For instance, the only thing that is contradictory in the KB is that jane threatening and jane stabbing, cannot both happen. In the final network (Figure 4.5), there are nodes for both threatening and stabbing. However, the relations of these nodes are such, that when one node becomes more likely (eg, we set evidence on it), the other node becomes less likely (as we can see in the figure below:

figure here of ΔP stabbing for increasing values of P threatening. If we know threatening is true, stabbing goes to 0, and vice-versa - and all the consequences of the threatening scenario also go to 0, so it just works. This means that several different scenarios can be combined into one network, and still be mutually exclusive. Both scenarios share some events (as some scenarios can do in real life), but the different scenarios themselves are not separated from each other by structure or construction (as is the case in Vlek 2015).

Do we need the constraint node? Or do we only use that if we have knowledge but no data? what. 11 in Vlek.

We can just straight-up draw a line between two events in order for them to become mutually exclusive, like we do with Jane stabs mark and jane threatens mark. Fenton et al advice against that - we draw a relation between two causes to ensure mutual exclusivity. Fenton doesn’t like this even though it satisfies the axioms because 1)

the parent cause becomes part of the causal pathway leading to the child cause, which means that you have to involve many more numbers (“meaningless columns”). Fortunately, this is not a problem in the automatically generated BNs, because the algorithm fixes this for us :). Not sure if an extra node is good for the computational complexity of BNs either. Uh. Anyway, even if this is not the case, there are lots of unnecessary numbers in Bayesian Networks anyway (show image here), because there’s many combinations of events that just do not happen (eg: mark dies but jane doesn’t have a knife, never happens, but is in the table anyway. Unnecessary complexity? Probably a rounding error!!). Anyway, that’s a thing. So it doesn’t really matter.

Uh. The second objection is that you have to arbitrarily decide which cause is the parent - which doesn’t make sense if you interpret the networks causally. Fortunately, there’s no causality in this part, its just frequencies so it doesn’t matter, we can just pick one and it’s fine. So turns out we don’t need the constraint node anyway :D

- Plausible - scenario should correspond to the modellers knowledge about the world. Support can help implausible scenarios to become plausible.

4.5 Important take-aways

In this section, I summarise some important take-aways, that will be necessary for the rest of this project. Things that we learn here, will be useful in evaluating the rest.

4.5.1 We don’t need the scenario node for mutual exclusivity

Both of Fenton’s objections for the direct node connection are not problems for us (yet) - the first one, specifying too many numbers, might become problematic in the future, and the second, about the interpretation of the arcs between nodes is not a problem, because we interpret these links as solely conditional, and not causal. Background information about causal relations between events might help us to order the network, so that it can become more efficient, but this should not affect the “reasoning” aspect - eg, even if the structure of the network is different, updating on evidence should produce the same results as before.

4.5.2 Resolution and number of runs

The accuracy of the Bayesian Network improves as the number of runs improves. This is an obvious fact. It has implications for the rest of the network. For example, if we run the network too few times (let’s say 10 times), then the network will not be able to estimate the probability of certain states - eg: the probability that Mark hits Jane can be calculated as: $0.7 \cdot 0.2 \cdot 0.03 \cdot 0.9 = 0.00379$, which means that there’s a 0.378% chance that it happens,

eg, if we run the simulation a 1000 times, Mark will hit Jane in 4 of them (or, actually, the sentence “Mark hits Jane” is true in 4 of them). If we go all the way down the chain of unreasonable facts to “Mark dies accidentally”, we get a probability of 0.0001134, which is 1 every 10.000 runs. If we run the simulation 10.000 times, that means that we cannot estimate probabilities that are smaller than 0.0001 - either they occur ‘accidentally’, eg, by chance, and the network estimates that they happen once every 10.000 runs, or they don’t happen within 10.000 runs, and the network estimates that these events never happen at all - when in fact, they might happen, but just once every 20.000 runs.

What are the practical implications of this? On this side, that we need enough data to get an accurate prediction. But, if we’re thinking about elicitation, this means that if you decide to add an event to a scenario, or a node to a network, you really have to think about the probability that you assign it - a probability of 1 in 10.000 would really mean that this event would only happen once every 10.000 situations (that you are considering - the reference class is never far away).

Chapter 5

A Simple Robbery - Spatial Simulations and their Bayesian Networks

5.1 Introduction

Here I talk about the robbery example with the evidence. Also model reference class explicitly just for fun here.

Chapter 6

The Grote Markt - investigating idioms with the Island Prior

6.1 Introduction

Here I talk about the Grote Markt and the more complex simulations that I created for it, including the island prior.

Chapter 7

Conclusion

7.1 Conclusion

Here I draw some conclusions.

7.2 Future Work

Testing or generating Bayesian Network idioms from simulations. The dream is to get “plug and play” Bayesian Network idioms - preconnected structures (perhaps even with some probabilities attached) that you can add evidence to and adapt and combine if necessary. Using simulations, we can test the granularity of these possible idioms, to simulate an crime at larger and smaller resolution (more or fewer events) to see how well the idioms can capture it.